

# CSC4005HW4: Heat-Distribution Simulation & Performance Analysis

118010202 / LIU Zhixuan

December 6, 2021

## 1 Abstract

The heat simulation generates a model in which there are four walls and a fireplace. The temperature of the wall is 20 degree , and the temperature of the fireplace is 100 degree. In this assignment, we make use of Jabobi iteration and Sor algorithm to compute the temperature inside the room and plot temperature contours at 5 degree intervals using Imgui.

In this assignment, I performed both parallel version and sequential version to draw the Mandelbrot Set using ImGui; moreover, for the parallel implementation, I use both MPI method and pthread method, OpenMp method, MPI&OpenMP method, and CUDA method to realize this simulation. At last, the corresponding experiments and some efficient analysis is performed on all the algorithms.



Figure 1: Illustration of Heat-Distribution Simulation

## 2 Introduction to Heat-Distribution Simulation and Sequential Realization

In this section, we briefly introduce Heat-Distribution Simulation and how to implement the sequential version of it using C++ and some complexity analysis of it.

Our goal is to initialize temperature of the total grid first, and then simulate the temperature state of the whole grid according to the temp passing algorithms. There are two algorithms that we use, the first one is Jacobi algorithm. Jacobi Iteration is invented to solve  $n \times n$  equations set. For a given matrix A, it can be decomposed into a diagonal component D, and the remainder R:  $A = D + R$ . The solution is then can be obtained iteratively via  $x_{k+1} = D^{-1}(b - Rx_k)$ , where  $x_k$  is the kth approximation. the formula is shown below:

$$x_i^{k+1} = \frac{1}{a_{ii}}(b_i - \sum_{j \neq i} a_{ij}x_j^k), i, j = 1, 2, \dots, n$$

The standard convergence condition is when the spectral radius of the iteration matrix is less than 1. A sufficient but not necessary condition for the method to converge is the matrix A is strictly or irreducibly diagonally dominant. Strict row diagonal dominance means that for each row, the absolute value of the diagonal term is greater than the sum of absolute values of other terms:  $|a_{ii}| < \sum_{j \neq i} |a_{ij}|$ .

Now, let's talk about the code implementation of this sequential algorithm using C++. In the following codes, i use ImGui to implement the drawing process.

Before the simulation, I receive some basic information to initialize the state and grid. These are the default initialization of the state and grid, which represent how many blocks we have, the size of the room, and some useful parameter.

---

```

1  int room_size = 300;
2  float block_size = 2;
3  int source_x = room_size / 2;
4  int source_y = room_size / 2;
5  float source_temp = 100;
6  float border_temp = 36;
7  float tolerance = 0.02;
8  float sor_constant = 4.0;
9  Algorithm algo = hdist::Algorithm::Jacobi;
```

---

Firstly, I define a buffer, which is called hdist::Grid, to store the information of the heat information. Its attributes are all vectors and integers, because they store the information of all heat, and the information of each block is determined by the index. Where data0 and data1 are different buffers, which store the information of the current stage, and the next stage. The reason why we use data0 and data1 is that we want the read and read to be separated.

---

```

1  struct Grid {
2      std::vector<double> data0, data1;
3      size_t current_buffer = 0;
4      size_t length;
5  }
```

---

Now let's move into the main loop. The first thing we do is to check whether the basic parameters have been changed or not in this iteration.

---

```

1  if (current_state.room_size != last_state.room_size) {
2      grid = hdist::Grid{
3          static_cast<size_t>(current_state.room_size),
4          current_state.border_temp,
5          current_state.source_temp,
6          static_cast<size_t>(current_state.source_x),
7          static_cast<size_t>(current_state.source_y)};
8      first = true;
9  }
```

---

We now go into the main function, the parameter finish is used to detect the finish time of the algorithm. If the program is finished, we will print out its finish time. And for the function implementation itself, we used two algorithms, one is Jacobi, another one is Sor algorithm.

---

```

1  if (!finished) {
2      finished = hdist::calculate(current_state, grid);
3      if (finished) end = std::chrono::high_resolution_clock::now();
4  }
5  bool calculate(const State &state, Grid &grid) {
6      bool stabilized = true;
7
8      switch (state.algo) {
9          case Algorithm::Jacobi:
10             for (size_t i = 0; i < state.room_size; ++i) {
11                 for (size_t j = 0; j < state.room_size; ++j) {
12                     auto result = update_single(i, j, grid, state);
13                     stabilized &= result.stable;
14                     grid[{alt, i, j}] = result.temp;
15                 }
16             }
17             grid.switch_buffer();
18             break;
19          case Algorithm::Sor:
20             for (auto k : {0, 1}) {
21                 for (size_t i = 0; i < state.room_size; i++) {
22                     for (size_t j = 0; j < state.room_size; j++) {
23                         if (k == ((i + j) & 1)) {
24                             auto result = update_single(i, j, grid, state);
25                             stabilized &= result.stable;
26                             grid[{alt, i, j}] = result.temp;
27                         } else {
28                             grid[{alt, i, j}] = grid[{i, j}];
29                         }
30                     }
31                 }
32                 grid.switch_buffer();
33             }
34         }
35         return stabilized;
36     };

```

---

Basically, in the calculate function, we go through each i,j pairs in the grid and update the temp information of it. We call the function which is update\_single, and we will select using different algorithms.

---

```

1  UpdateResult update_single(size_t i, size_t j, Grid &grid, const State &state) {
2      UpdateResult result{};
3      if (i == 0 || j == 0 || i == state.room_size - 1 || j == state.room_size - 1) {
4          result.temp = state.border_temp;
5      } else if (i == state.source_x && j == state.source_y) {
6          result.temp = state.source_temp;
7      } else {
8          auto sum = (grid[{i + 1, j}] + grid[{i - 1, j}] + grid[{i, j + 1}] + grid[{i, j - 1}]);
9          switch (state.algo) {
10              case Algorithm::Jacobi:
11                  result.temp = 0.25 * sum;

```

```

12         break;
13     case Algorithm::Sor:
14         result.temp = grid[{i, j}] + (1.0 / state.sor_constant) * (sum - 4.0
15             * grid[{i, j}]);
16         break;
17     }
18     result.stable = std::fabs(grid[{i, j}] - result.temp) < state.tolerance;
19     return result;
20 }

```

---

After updating the information write them into the buffer, next we will completely update the information of the target  $i$  one by one for this iteration using the information we calculated above using the function above.

After update all the information in this iteration, we now comes to the final part, which is drawing all the new updates on the ImGui.

Basically, the main idea to draw the things is to go through each body, which is represented by  $i$  in each iteration, and then draw it on the ImGui.

```

1  const ImVec2 p = ImGui::GetCursorScreenPos();
2  float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
3  for (size_t i = 0; i < current_state.room_size; ++i) {
4      for (size_t j = 0; j < current_state.room_size; ++j) {
5          auto temp = grid[{i, j}];
6          auto color = temp_to_color(temp);
7          draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x + current_state.block_size,
8              y + current_state.block_size), color);
9          y += current_state.block_size;
10     }
11     x += current_state.block_size;
12     y = p.y + current_state.block_size;
13 }
14 ImGui::End();

```

---

Now, let's evaluate the complexity of sequential Heat-Distribution simulation computation algorithm. The computational complexity :  $O(NM)$  where,  $N$  and  $M$  are the room size (length). Because we will go through each block for the computation and the total number of blocks are  $NM$ . In our case, the computational complexity is  $O(N^2)$ .

### 3 Parallel Heat-Distribution Simulation Computation using MPI

It is obvious that sequential Heat-Distribution simulation computation does not have a very good performance in time complexity. However, this algorithm can be much more efficient if it is distributed computed, therefore, the computational time will be saved. The first idea of doing this is to use MPI method to spread up the process, and the main idea is as follows:

In the MPI algorithm, firstly, the assignment will be given and initialized in the master process. Then the master process will distribute the information in the slave processes, each slave process is equally assigned to calculate some part of the  $n$ -bodies, and only up date its information. After calculate and update their own parts, they will sent the result to the master process. And the master process will write the information in the buffer "grid", and then using ImGui to visualize the body moving.

The process for this distributed computing is shown in the figure 2 below:




Figure 2: Process of Parallel computation using MPI

Now let's look at how to implement this process using C++.

Similar to the sequential process, we first initiate a structure called pool, which stored the information of all the bodies in the master process, and initialize some necessary parameters.

Next, in the MASTER process, the MASTER will send the parameter information to all the slave process by using iteration:

---

```

1 for(int i=0; i<slave_size; i++){
2     MPI_Send(&(para_Size[0]), 2, MPI_INT, (i+1), (i+1)*10, MPI_COMM_WORLD);
3     MPI_Send(&(para_State_int[0]), 4, MPI_INT, (i+1), (i+1)*10+1, MPI_COMM_WORLD);
4     MPI_Send(&(para_State_float[0]), 5, MPI_FLOAT, (i+1), (i+1)*10+2,
5         MPI_COMM_WORLD);
6     MPI_Send(&(g_data0[0]), m_length, MPI_DOUBLE, (i+1), (i+1)*10+3, MPI_COMM_WORLD
7         );
8     MPI_Send(&(g_data1[0]), m_length, MPI_DOUBLE, (i+1), (i+1)*10+4, MPI_COMM_WORLD
9         );
10    MPI_Send(&(para_Grid[0]), 2, MPI_INT, (i+1), (i+1)*10+5, MPI_COMM_WORLD);
11 }
12 MPI_Barrier(MPI_COMM_WORLD);

```

---

Then, in the slave process, where all the computations are distributed, slave first receive all the information sent by MASTER process:

---

```

1 MPI_Recv(&(s_data0[0]), s_length, MPI_DOUBLE, 0, rank*10+3, MPI_COMM_WORLD, &status);
2 MPI_Recv(&(s_data1[0]), s_length, MPI_DOUBLE, 0, rank*10+4, MPI_COMM_WORLD, &status);

```

---

---

```
3 MPI_Recv(&(s_para_Grid[0]), 2, MPI_INT,0,rank*10+5,MPI_COMM_WORLD, &status);
```

---

Then, only calculate part of the bodies and store the body information in the local buffer, the culculation process is the same as the sequential process:

---

```
1 bool s_finished = false;
2 s_finished = hdist::calculate(s_current_state, s_grid, s_local_size, s_room_size,
    rank);
```

---

After calculate their own part, the slave will send the buffer to the MASTER process, and the slave process finishes its job.

---

```
1 MPI_Isend(&s_finished_int, 1, MPI_INT, 0, 1, MPI_COMM_WORLD, &request);
2 MPI_Isend(&(s_grid.current_buffer), 1, MPI_INT, 0, 2, MPI_COMM_WORLD, &request);
3 MPI_Isend(&(s_data0[0]), s_length, MPI_DOUBLE, 0, 3, MPI_COMM_WORLD, &request);
4 MPI_Isend(&(s_data1[0]), s_length, MPI_DOUBLE, 0, 4, MPI_COMM_WORLD, &request);
```

---

Then, the MASTER process receive the buffer from the slave processes, and write the information to the grid:

---

```
1 MPI_Recv(&(g_data0[0]), m_length, MPI_DOUBLE,MPI_ANY_SOURCE, 3, MPI_COMM_WORLD,&
    status);
2 slave_rank = status.MPI_SOURCE;
3 for (int j=0;j<local_size*room_size; j++){
4     if(((slave_rank-1)*local_size*room_size + j) < room_size*room_size){
5         grid.data0[(slave_rank-1)*local_size*room_size + j] = g_data0[(slave_rank
            -1)*local_size*room_size + j];
6     }
7 }
8
9 MPI_Recv(&(g_data1[0]), m_length, MPI_DOUBLE,MPI_ANY_SOURCE, 4, MPI_COMM_WORLD,&
    status);
10 slave_rank = status.MPI_SOURCE;
11 for (int j=0;j<local_size*room_size; j++){
12     if(((slave_rank-1)*local_size*room_size + j) < room_size*room_size){
13         grid.data1[(slave_rank-1)*local_size*room_size + j] = g_data1[(slave_rank
            -1)*local_size*room_size + j];
14     }
15 }
```

---

The last part is for canvas to draw the information using ImGui:

---

```
1 for (size_t i = 0; i < pool.size(); ++i) {
2     auto body = pool.get_body(i);
3     auto x = p.x + static_cast<float>(body.get_x());
4     auto y = p.y + static_cast<float>(body.get_y());
5     draw_list->AddCircleFilled(ImVec2(x, y), radius, ImColor{color});
6 }
```

---

This is pretty much about the concise implementation of the parallel Heat-Distribution simulation Computation using MPI.

As for the computational complexity, since the whole mathematical computational algorithm does not change, the only change is that the whole process is distributed to p slave processes to compute, therefore, the computational complexity can be written as in our cases are  $O(\frac{N^2}{p})$ .

## 4 Parallel Heat-Distribution Simulation Computation using Pthread

It is obvious that sequential computation does not have a very good performance in time complexity. However, this algorithm can be much more efficient if it is distributed computed, therefore, the computational time will be saved. The second idea of doing this is to use Pthread method to speed up the process, and the main idea is as follows:

In the Pthread algorithm, firstly, the assignment will be given and initialized in the main function. Then the main function will fork several child processes, and will distribute the information in the child processes, each child process is equally assigned to on part of the whole picture, and their aim is to draw their own part of the picture. Each child process will write the information in the buffer "grid". After writing the information, the each thread will join together, and then using ImGui to visualize the drawing.

The process for this distributed computing is shown in the figure 3 below:




Figure 3: Process of Parallel computation using Pthread

Now let's look at how to implement this Pthread process using C++.

Similar to the sequential process, we first initiate a structure called pool, which stored the information of the drawing based on the pixel in the master process, and initialize some necessary parameters.

Next, in the main function process, we will create a space to store some threads, based on the given threads number:

---

```
1 pthread_t threads[THREADS_NUM];
```

---

Then, we also create a structure called argument, which will save all the information that each thread's function will receive, and create and send the information to thread:

---

```
1 struct Arguments{
2     hdist::State * state;
3     hdist::Grid * grid;
4     bool * stable_flag;
5     int pid;
```

---

```

6     int local_bodies;
7     int bodies;
8 };
9
10 for(int m=0; m<THREADS_NUM; m++){
11     pthread_create(&threads[m], nullptr, local_process, new Arguments{
12         .state = &current_state,
13         .grid = &grid,
14         .stable_flag = &stable_flag,
15         .pid = m,
16         .local_bodies = local_size,
17         .bodies = current_state.room_size
18     });
19 }
20 }

```

---

Then, in the child process, firstly, they receive the information, and then only calculate part of the image and store the image information in the pool. To let them all write to the pool correctly, I passed the pointer of the canvas so that each thread can find the address to write to:

---

```

1 void *local_process(void *arg_ptr){
2     auto arguments = static_cast<Arguments *>(arg_ptr);
3     // PTHREAD_MUTEX_LOCK(&MUTEX_P);
4     bool local_stable;
5     local_stable = hdist::calculate(*(arguments->state), *(arguments->grid),
6                                     arguments->pid+1, arguments->local_bodies,
7                                     arguments->bodies);
8     // PTHREAD_MUTEX_UNLOCK(&MUTEX_P);
9     *(arguments->stable_flag) &= local_stable;
10    delete arguments;
11    return nullptr;
12 }

```

---

After calculate their own part, the all the threads will join together, and the whole process ends.

---

```

1 for(auto & n : threads){
2     pthread_join(n, nullptr);
3 }
4 pthread_attr_destroy(&attr);

```

---

The last part is for pool to draw the information using ImGui:

---

```

1 const ImVec2 p = ImGui::GetCursorScreenPos();
2 float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
3 for (size_t i = 0; i < current_state.room_size; ++i) {
4     for (size_t j = 0; j < current_state.room_size; ++j) {
5         auto temp = grid[{i, j}];
6         auto color = temp_to_color(temp);
7         draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x + current_state.block_size,
8                                                         y + current_state.block_size), color);
9         y += current_state.block_size;
10    }
11    x += current_state.block_size;
12    y = p.y + current_state.block_size;
13 }
14 ImGui::End();

```

---



This is pretty much about the concise implementation of the parallel N-body Simulation Computation using Pthread.

As for the computational complexity, since the whole mathematical computational algorithm does not change, the only change is that the whole process is distributed to  $c$  child processes to compute, therefore, the computational complexity can be written as in our cases are  $O(\frac{N^2}{c})$ .

## 5 Parallel Heat-Distribution Simulation Computation using OpenMP

It is obvious that sequential computation does not have a very good performance in time complexity. However, this algorithm can be much more efficient if it is distributed computed, therefore, the computational time will be saved. The third idea of doing this is to use OpenMP method to speed up the process, and the main idea is as follows:

In the OpenMP algorithm, similar to the Pthread algorithm, the only difference is the way it implement. Instead of fork child process in using pthread, it is used in this way.

---

```

1 {
2  if (!finished) {
3      omp_set_num_threads(THREADS_NUM);
4      finished = hdist::calculate(current_state, grid);
5      if (finished) end = std::chrono::high_resolution_clock::now();
6  }
7
8  bool calculate(const State &state, Grid &grid) {
9      bool stabilized = true;
10     int i;
11     {
12         switch (state.algo) {
13             case Algorithm::Jacobi:
14                 #pragma omp for schedule(static)
15                 for (i = 0; i < state.room_size; ++i) {
16                     for (int j = 0; j < state.room_size; ++j) {
17                         auto result = update_single(i, j, grid, state);
18                         stabilized &= result.stable;
19                         grid[{alt, i, j}] = result.temp;
20                     }
21                 }
22                 grid.switch_buffer();
23                 break;
24             case Algorithm::Sor:
25                 for (auto k : {0, 1}) {
26                     #pragma omp for schedule(static)
27                     for (i = 0; i < state.room_size; i++) {
28                         for (int j = 0; j < state.room_size; j++) {
29                             if (k == ((i + j) & 1)) {
30                                 auto result = update_single(i, j, grid, state);
31                                 stabilized &= result.stable;
32                                 grid[{alt, i, j}] = result.temp;
33                             } else {
34                                 grid[{alt, i, j}] = grid[{i, j}];
35                             }
36                         }
37                     }

```

```

38         grid.switch_buffer();
39     }
40 }
41 }
42 return stabilized;
43 };

```

---

The process for this distributed computing is shown in the figure 4 below:

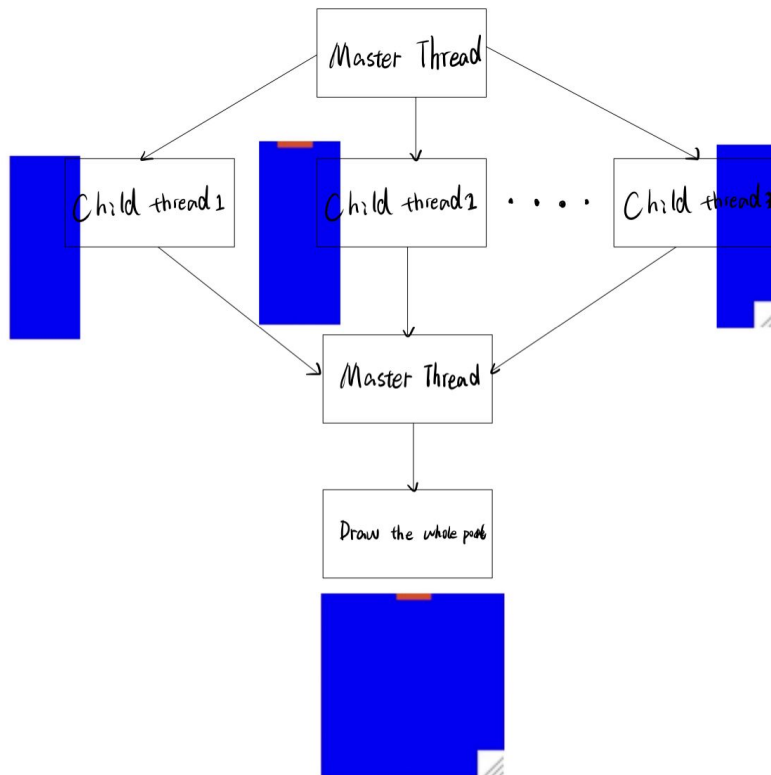


Figure 4: Process of Parallel computation using OpenMP

## 6 Bonus: Parallel N-body Simulation Computation using OpenMP and MPI

It is obvious that sequential computation does not have a very good performance in time complexity. However, this algorithm can be much more efficient if it is distributed computed, therefore, the computational time will be saved. The third idea of doing this is to use MPI and OpenMP method to spread up the process, and the main idea is as follows:

This algorithm is based on the MPI. After the MASTER process distribute jobs to the SLAVE process, in each SLAVE process, it can call the OpenMP to generate more threads in each SLAVE process. Therefore there will be more threads to implement the job, each process get lesser jobs. The differences in codes are shown below, where the baseline is the MPI algorithm process.

---

```

1 {

```

```

2  if (!finished) {
3      omp_set_num_threads(THREADS_NUM);
4      finished = hdist::calculate(current_state, grid);
5      if (finished) end = std::chrono::high_resolution_clock::now();
6  }
7
8  bool calculate(const State &state, Grid &grid) {
9      bool stabilized = true;
10     int i;
11     {
12         switch (state.algo) {
13             case Algorithm::Jacobi:
14                 #pragma omp for schedule(static)
15                 for (i = 0; i < state.room_size; ++i) {
16                     for (int j = 0; j < state.room_size; ++j) {
17                         auto result = update_single(i, j, grid, state);
18                         stabilized &= result.stable;
19                         grid[{alt, i, j}] = result.temp;
20                     }
21                 }
22                 grid.switch_buffer();
23                 break;
24             case Algorithm::Sor:
25                 for (auto k : {0, 1}) {
26                     #pragma omp for schedule(static)
27                     for (i = 0; i < state.room_size; i++) {
28                         for (int j = 0; j < state.room_size; j++) {
29                             if (k == ((i + j) & 1)) {
30                                 auto result = update_single(i, j, grid, state);
31                                 stabilized &= result.stable;
32                                 grid[{alt, i, j}] = result.temp;
33                             } else {
34                                 grid[{alt, i, j}] = grid[{i, j}];
35                             }
36                         }
37                     }
38                     grid.switch_buffer();
39                 }
40             }
41     }
42     return stabilized;
43 };

```

---

The process for this distributed computing is shown in the figure 5 below:

## 7 Parallel Heat-Distribution Simulation Computation using Cuda

It is obvious that sequential computation does not have a very good performance in time complexity. However, this algorithm can be much more efficient if it is distributed computed, therefore, the computational time will be saved. The second idea of doing this is to use Cuda method to spread up the process, and the main idea is as follows:

In the Cuda algorithm, firstly, the assignment will be given and initialized in the main function. Then the main function will fork several thread processes, and will distribute the information in the thread processes, each thread process is equally assigned to on

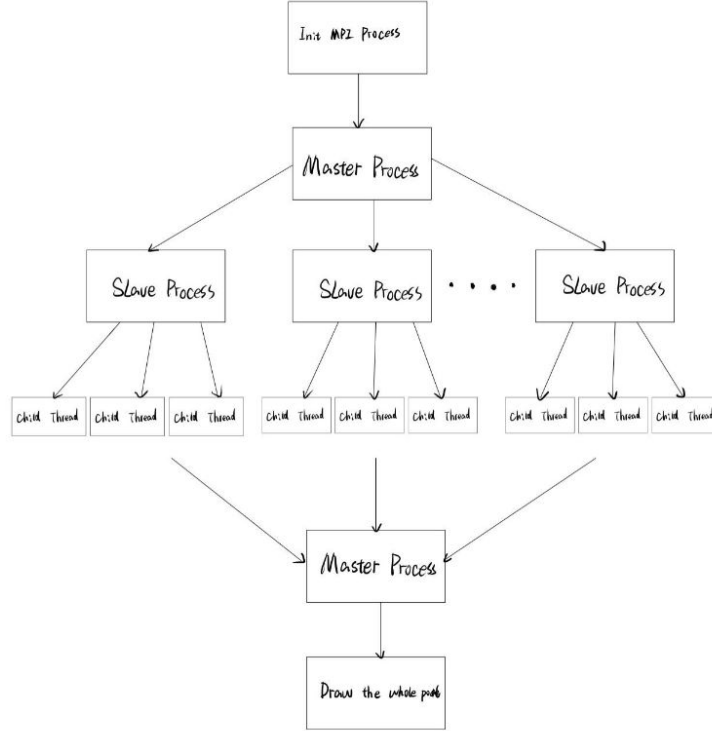


Figure 5: Process of Parallel computation using MPI and OpenMP

part of the  $n$  bodies, and their aim is to draw their own part of the bodies. Each thread process will write the information in the buffer "grid". After writing the information, the each thread will join together, and then using ImGui to visualize the drawing.

The process for this distributed computing is shown in the figure 6 below:

The major differences are `std::vector` cannot be used in the cuda process, so we mainly swift to array base implementation. The second thing is that the array initialization is in the main function, but cuda devices need to have access to it. Therefore, we use `Cudamemcpy` to let cuda devices to get access to the array and change on it.

---

```

1 double * Hostdata0 = new double[room_size_C*room_size_C];
2 double * Hostdata1 = new double[room_size_C*room_size_C];
3 bool *Hostdevicestable = new bool;
4
5 double *data0;
6 cudaMalloc(&data0, sizeof(double) * room_size_C*room_size_C);
7 double *data1;
8 cudaMalloc(&data1, sizeof(double) * room_size_C*room_size_C);
9 bool *devicestable;
10 cudaMalloc(&devicestable, sizeof(bool));

```

---

This is pretty much about the concise implementation of the parallel Heat-Distribution Simulation Set Computation using Cuda.

As for the computational complexity, since the whole mathematical computational algorithm does not change, the only change is that the whole process is distributed to  $t$  thread processes to compute, therefore, the computational complexity can be written as in our cases are  $O(\frac{N^2}{t})$ .

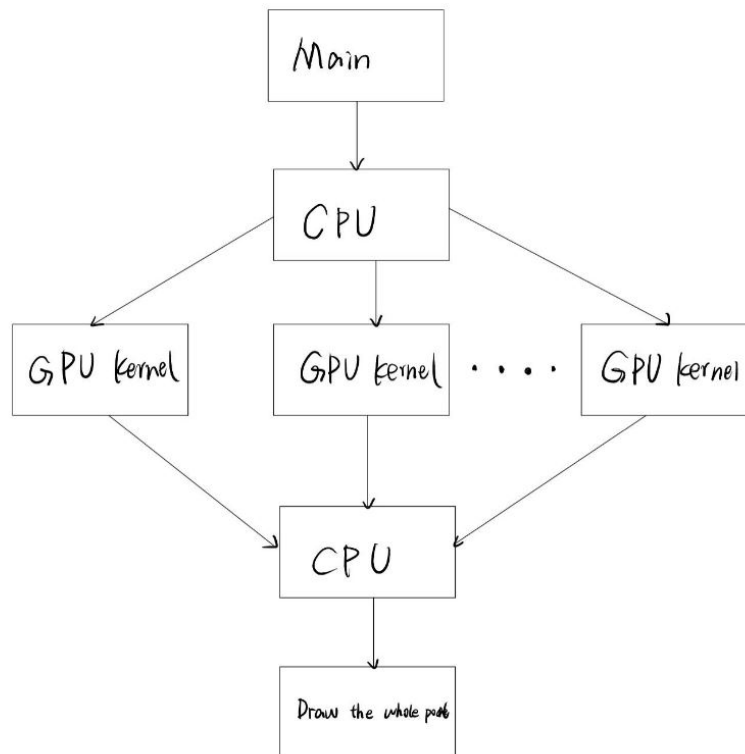


Figure 6: Process of Parallel computation using Cuda

## 8 Compile and Run

I use the template on the blackboard to write the program, which use ImGui to display the image. So here are some command to run my code.

IMPORTANT: Please put my codes under file src and use the file name: "main.cpp". And put the "hdist.hpp" file to the under the /include/hdist folder, and also change the CMakeList.txt in each version of algorithm.

### 8.1 Build

For sequential, MPI, pthread, OpenMp, MPI + OpenMP, please do the following to build the file:

---

```

1 $ rm -rf build
2 $ mkdir build
3 $ cd build
4 $ cmake ..
5 $ make
  
```

---

as for the cuda version:

---

```

1 $ cd csc4005-imgui
2 $ mkdir build && cd build
3 $ source scl_source enable devtoolset-10
4 $ CC=gcc CXX=g++ cmake ..
5 $ make -j12
  
```

---

## 8.2 Run Sequential Program

To run my sequential codes, the sbatch shell codes are as follows:

---

```
1 #!/bin/bash
2 #SBATCH -J Zhixuan
3 #SBATCH -o sequential.out
4 #SBATCH -N 1
5 #SBATCH -n 1
6 #SBATCH -t 10
7
8 xvfb-run -a mpirun -np 1 csc4005_imgui
```

---

## 8.3 Run MPI Program

To run my MPI codes, the sbatch shell codes are as follows:

---

```
1 #!/bin/bash
2 #SBATCH -J Zhixuan
3 #SBATCH -o MPI_N1n4.out
4 #SBATCH -N 1
5 #SBATCH -n 4
6 #SBATCH -t 10
7
8 xvfb-run -a mpirun -np 4 csc4005_imgui
```

---

## 8.4 Run Pthread Program

To run my Pthread codes, the sbatch shell codes are as follows:

---

```
1 #!/bin/bash
2 #SBATCH -J Zhixuan
3 #SBATCH -o Pthread.out
4 #SBATCH -N 1
5 #SBATCH -n 1
6 #SBATCH -t 10
7
8 xvfb-run -a mpirun -np 1 csc4005_imgui
```

---

To test different thread, please change the thread number in the first line of the code in main.cpp, where you can modify the predefined THREADS\_NUM.

## 8.5 Run OpenMP Program

To run my OpenMP codes, the sbatch shell codes are as follows:

---

```
1 #!/bin/bash
2 #SBATCH -J Zhixuan
3 #SBATCH -o OpenMP.out
4 #SBATCH -N 1
5 #SBATCH -n 1
6 #SBATCH -t 10
7
8 xvfb-run -a mpirun -np 1 csc4005_imgui 4
```

---

To test different thread, please change the thread number mentioned above, the default number will be 4.

## 8.6 Run MPI+OpenMP Program

To run my MPI+OpenMP codes, the sbatch shell codes are as follows:

---

```
1 #!/bin/bash
2 #SBATCH -J Zhixuan
3 #SBATCH -o MPI_N1n4.out
4 #SBATCH -N 1
5 #SBATCH -n 4
6 #SBATCH -t 10
7
8 xvfb-run -a mpirun -np 4 csc4005_imgui
```

---

To test different thread, please change the thread number in the first line of the code, where you can modify the predefined THREADS\_NUM in the main.cpp.

## 8.7 Run Cuda Program

To run my OpenMP codes, the sbatch shell codes are as follows:

---

```
1 #!/bin/bash
2 #SBATCH -J Zhixuan
3 #SBATCH -o Cuda.out
4 #SBATCH -N 1
5 #SBATCH -n 1
6 #SBATCH -t 10
7
8 xvfb-run -a mpirun -np 1 csc4005_imgui 4
```

---

To test different cuda thread, please change the thread number mentioned above, the default number will be 4.

## 8.8 Sample Outcomes

Here are some sample outcomes for both sequential, mpi, and pthread, openmp, cuda algorithms.

# 9 Analysis

## 9.1 Comparison between Sequential, MPI, Pthread, MPI\_Pthread and Cuda

As the figure 8 shown above, when we fix the number of processes used by the program for MPI and Pthread (here we are talking about the fixed number, regardless of the number of masters in the mpi program, because when the MPI is executing the program, the MASTER program actually does not perform calculations, and It is to allocate tasks. In order to make the comparison more convincing, the number of processes mentioned here is slave process and child process). Adjust the job size of the program. The graph we get reflects the growth of time. In this figure, we can see that in the 6 programs, time has a growth trend similar to a quadratic function with the increase of job size. This is in line with our previous analysis that the computational complexity of the three algorithms is quadratic.

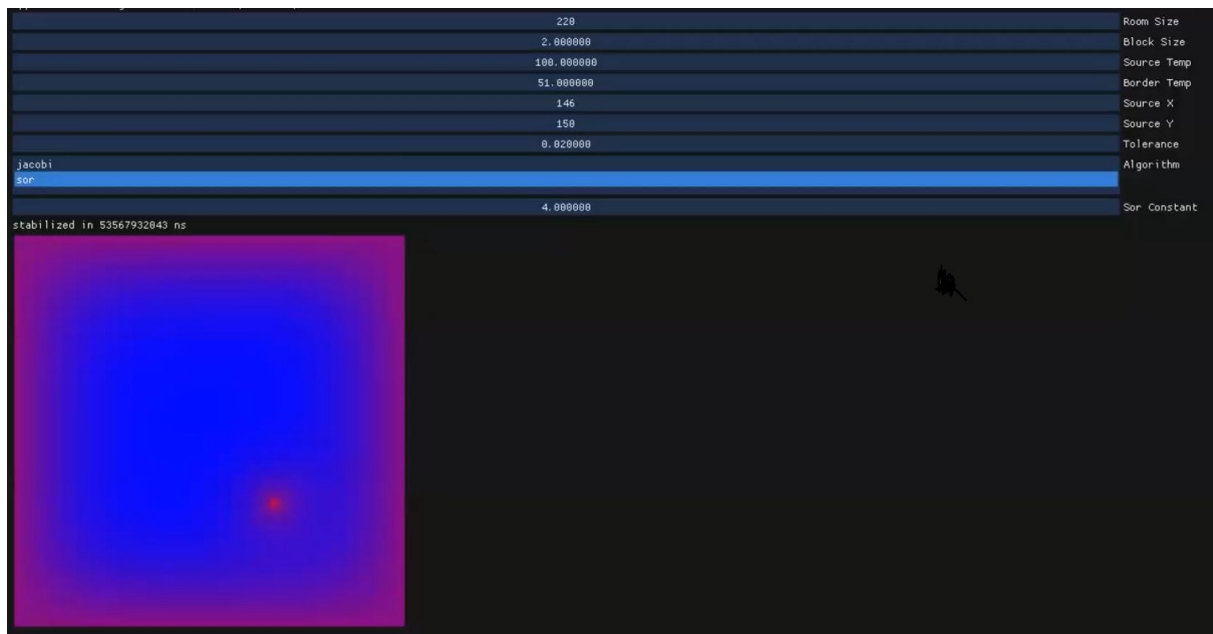


Figure 7: Sample GUI outcomes

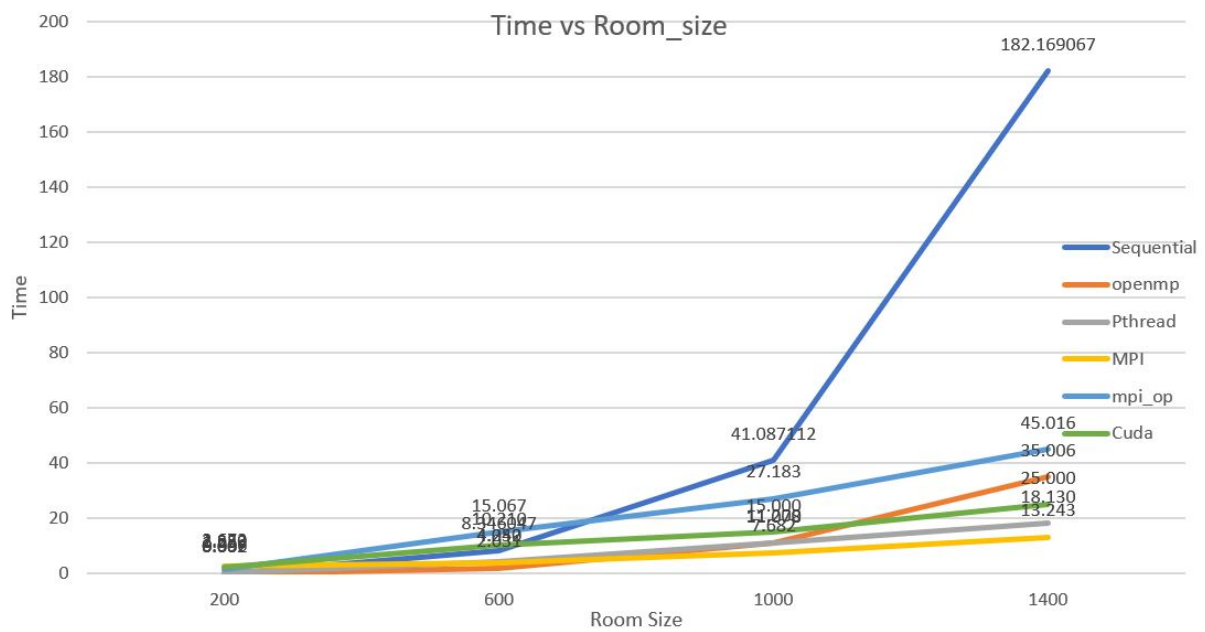


Figure 8: Comparison between all the algorithms



Firstly, since all the programs take shorter time than the sequential program, it indicates that our program indeed has the performance improvement. At the same time, we found that mpi and pthread programs are shorter in time than sequential programs, indicating that our distributed design can accelerate computing efficiency the most. We found that when the process number is 12, distributed programs accelerate the program more than when the process number is 12. This observation is consistent with the computational complexity of distributed programs mentioned earlier, and we will analyze the specific multiples later.

## 9.2 Comparison focusing on one algorithm

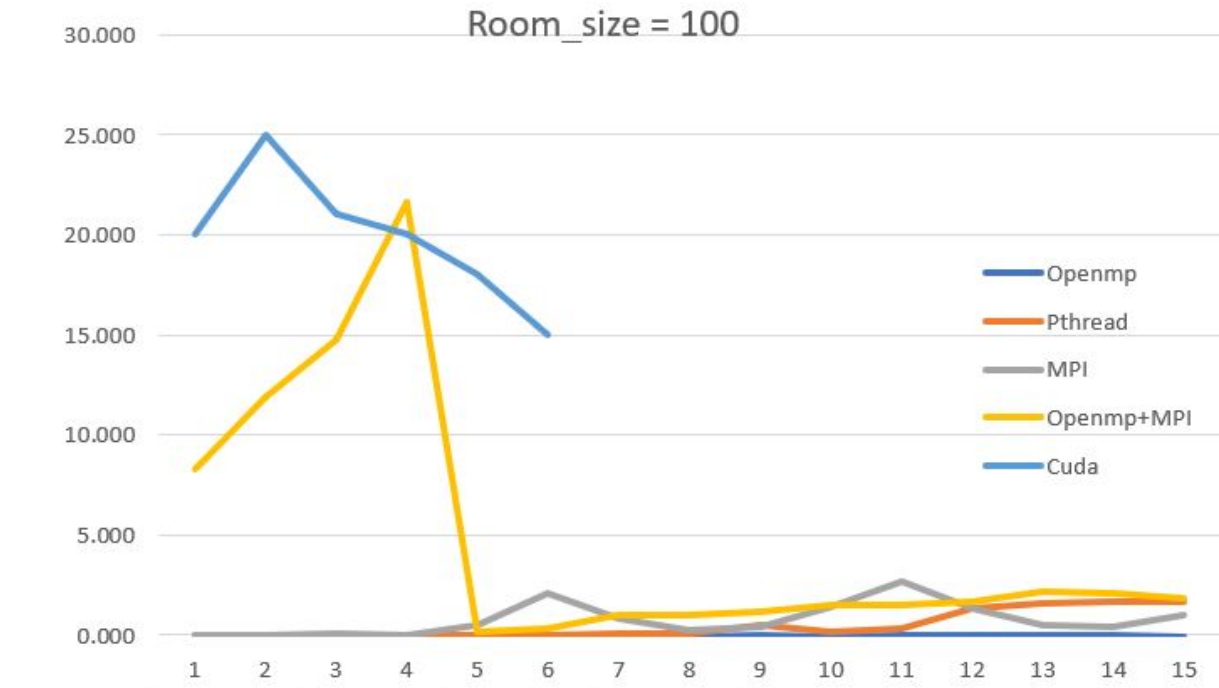


Figure 9: Comparison between all the algorithms: Time vs Process numbers when bodies number is fixed to be 100

The above two pictures show the relationship between the time spent by 6 algorithms and the process number when the job size is fixed. For the fairness of comparison, this type of process number refers to the slave process of mpi and the child process of pthread.

The first figure 9 shows that when the body number is fixed at 100, how the time those algorithm will spend when the process numbers is increasing. However, it is anomalous that when the number of processors increases, the family uses them longer. This may be because, because the number of bodies is too small, each processor cannot allocate many bodies, so the locked time does not depend on the time used for calculation, but depends on the allocated processor, read Time spent writing and transferring memory. For example, the time used by cuda in this process is relatively old. The reason may be that I have a lot of Cudamecpy in my cuda program. This leads to a waste of time, so cuda will take a long time.

Now since the job size is much bigger, we can observe some useful information. We found that when the number of processes increases, the time has a slight downward trend as shown in figure 10. This trend is consistent with the computational complexity of the 5




Figure 10: Comparison between all the algorithms: Time vs Process numbers when bodies number is fixed to be 500

previously analyzed. When the number of process numbers is larger, the amount of work allocated to each process will also become less, so the time taken to complete a single computing task will also become less. This effectively shows that distributed computing can reasonably increase the time used. And this tendency is more obvious when the body size is more than 1000, as shown in figure

At the same time, we also found that although the time used by mpi is very close to that of pthread, it still takes more time than pthread within the tolerance range. This point may be the same as the reason we analyzed in the previous section. Because in pthread, all child processes will be rewritten directly on the canvas, they do not need a local buffer to store data. But in my mpi algorithm, I use the local buffer to store the data of the slave processes and pass this data to the MASTER process. At the same time, the MATER process also needs to write the received signal into the canvas. This process is very Time-consuming.

Moreover, I also observed that cuda has a very outstanding performance when body size is increasing. Although its performance may not be super good at the beginning when the processor number is small, but it is super cool when the processor number is increasing.

## 10 Conclusion and Future work

This focuses on the subject that how to write static scheduling program and write a pthread program and most importantly, OpenMP program and integrate the OpenMP with the MPI method. In this project, the N-body Simulation image is displayed by using Sequential, MPI, OpenMP, Pthread, MPI-OpenMP Hybrid methods, and cuda implementation. The above performance analysis clarifies why the MPI method will show a downward trend of performance and both of the Pthread and OpenMP have better




Figure 11: Comparison between all the algorithms: Time vs Process numbers when bodies number is fixed to be 1000

performance than MPI when the body size is small and medium. When the body size grows, both of the MPI method and Pthread method shows a general downward trend of running time. The MPI method has a improvement in performance. And cuda performe the best when the job size is large and the device number increases.

Moreover, I also find that in the same condition, my implimentation of pthread is actually much faster than MPI. This is due to the face that in my practice, pthread directly write to the buffer which store the information of generated image while in MPI I need to use local buffer and then send the information to the master process, then the master process writes to the image buffer. This is actually time consuming point. And this point can be further improved.

Moreover, due to the resource limitation, some subtle and comprehensive experiment should be done in the further situation.